

Lab 8

Lindsey China (A17023629)

Table of contents

Importing and Formatting Data	1
PCA on wisc.data	5
Interpreting PCA Results	7
Variance Explained	9
Communicating PCA Results	11
Hierarchical Clustering	12
Selecting number of clusters	13
Using different methods	14
K-means clustering	17
Combining Methods	17
Clustering on PCA results	17
Sensitivity and specificity	21
Prediction	22

Importing and Formatting Data

Need to import the Wisconsin cancer data for this project, after loading it into the directory create a data frame to be used in the code:

```
fna.data <- "WisconsinCancer.csv"

wisc.df <- read.csv(fna.data, row.names=1)
```

```
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217

	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137

	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41
84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75

	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663

84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302	0.11890			
842517	0.08902			
84300903	0.08758			
84348301	0.17300			
84358402	0.07678			
843786	0.12440			

The first column won't be used here, it is a pathologist diagnosis and is basically the answer to if cells are malignant or benign. Reformat the dataset to remove the first column:

```
wisc.data <- wisc.df[, -1]
```

Need to create a new vector, "Diagnosis", with the data from the diagnosis column of the original set to check our results later:

```
# To create a factor, need to use the tidyverse package
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2    3.5.0      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.1
v purrr      1.0.2
```

```
-- Conflicts ----- tidyverse_conflicts() --
```

```
x dplyr::filter() masks stats::filter()
```

```
x dplyr::lag()     masks stats::lag()
```

```
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
# Specify the diagnosis levels/variables present in the list
diagnosis_levels <- c("B", "M")
```

```
# Write the factor
diagnosis <- factor(wisc.df$diagnosis, levels=diagnosis_levels)
head(diagnosis)
```

```
[1] M M M M M M
Levels: B M
```

Q1. How many observations are in the dataset?

```
dim(wisc.df)
```

```
[1] 569 31
```

569 Observations of 31 variables

Q2. How many of the observations have a malignant diagnosis?

```
table(wisc.df$diagnosis)
```

```
 B   M
357 212
```

212 Observations are malignant

Q3. How many variables/features in the data are suffixed with __mean?

```
length((grep("__mean", names(wisc.df))))
```

```
[1] 10
```

10 Variables contain __mean

PCA on wisc.data

To determine if data needs to be scaled, check the column means and standard deviation:

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03

symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

Execute the PCA:

```
wisc.pr <- prcomp(wisc.data, scale. = TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From the results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27% of the variance is captured by PC1.

Q5. How many principle components are required to describe at least 70% of the original variance in the data?

3 PCs are required to capture 70% of the variance.


Q6. How many principal components are required to describe at least 90% of the original variance in the data?

7 PCs are required to capture 90% of the variance.

Interpreting PCA Results

Use a biplot to visualize the results of the PCA:

```
biplot(wisc.pr)
```

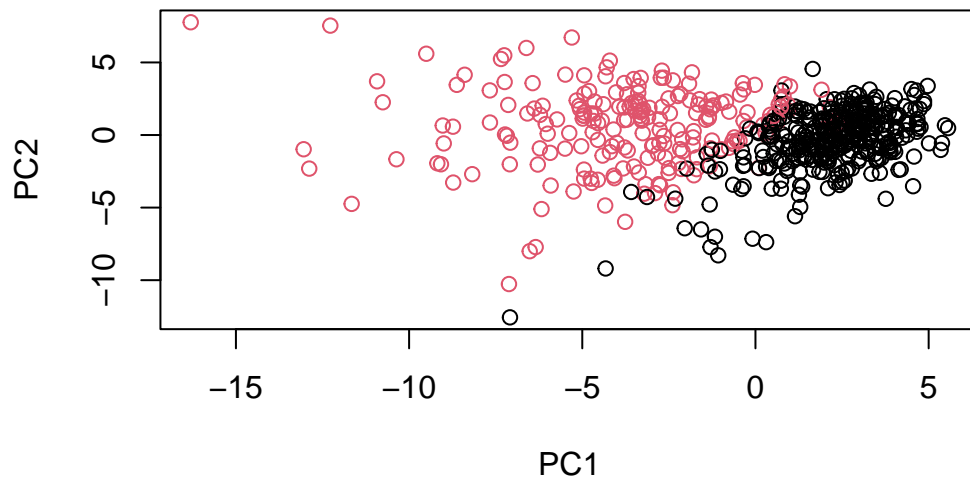


Q7. What stands out about this plot? Is it easy to understand? Why?

It is very difficult to understand, the data points are all on top of each other and it is difficult to see a pattern or meaning in the graph.

Try a scatter plot colored by the first two PCAs:


```
plot(wisc.pr$x,col=diagnosis,  
      xlab="PC1",ylab="PC2")
```



We can use ggplot2 to make an even better representation of this data:

```
# Create a data frame for ggplot to use
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

library(ggplot2)
ggplot(df)+
  aes(PC1, PC2, col=diagnosis)+
  geom_point()
```

Variance Explained

Calculate the variance from the PCA:

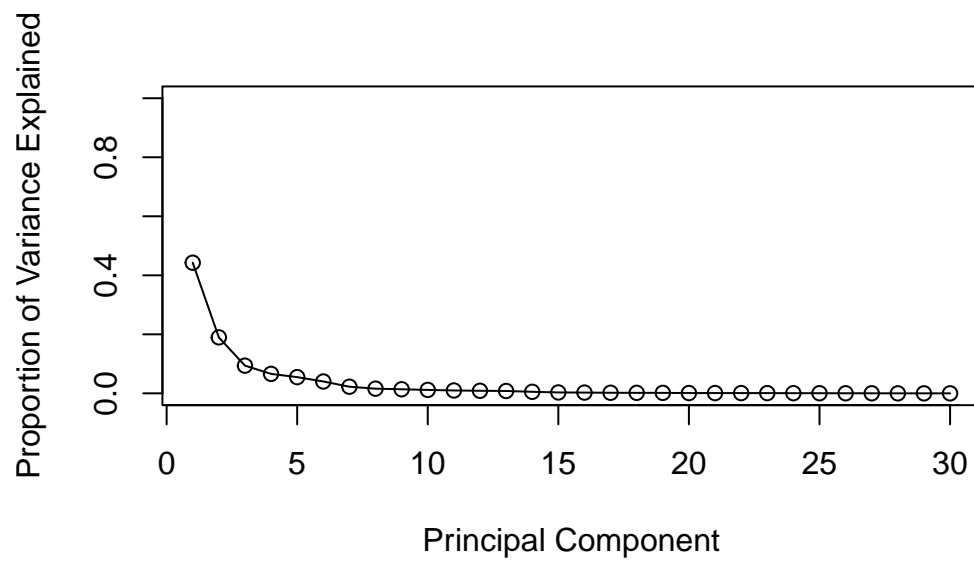
```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

Find the variance explained by each PC:


```
# Variance explained by each principal component, stored as pve
pve <- pr.var/sum(pr.var)

# Plot pve
plot(pve, xlab="Principal Component",ylab="Proportion of Variance Explained",ylim=c(0,1),t
```



Alternative plot of the same data with a data driven y-axis:

```
barplot(pve,ylab="Percent of Variance Explained",names.arg=paste0("PC",1:length(pve)),las=
axis(2,at=pve,labels=round(pve,2)*100)
```



Communicating PCA Results

Q9. For the first PC, what is the component of the loading vector for the feature `concave.points_mean`?

-0.2608538

```
head(wisc.pr$rotation[,1],8)
```

<code>radius_mean</code>	<code>texture_mean</code>	<code>perimeter_mean</code>	<code>area_mean</code>
-0.2189024	-0.1037246	-0.2275373	-0.2209950
<code>smoothness_mean</code>	<code>compactness_mean</code>	<code>concavity_mean</code>	<code>concave.points_mean</code>
-0.1425897	-0.2392854	-0.2584005	-0.2608538

Q10. What is the minimum number of principal components required to explain 80% of the variance in the data?

At least 5 PCs are required to explain 80% of the data

Hierarchical Clustering

Scale the data to prepare it for hierarchical clustering:

```
data.scaled <- scale(wisc.data)
```

Calculate the euclidean distance between all pairs of observations:


```
data_dist <- dist(data.scaled, method="euclidean")
```

Create a hierarchical clustering model using complete linkage, manually specify the method to `hclust()` and assign the results to `wisc.hclust`:

```
wisc.hclust <- hclust(data_dist, method="complete")
```

Q11. Using the `plot()` and `abline()` function, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)  
abline(h=19, col="red", lty=2)
```



At a height of 19 the clustering model is broken up into 4 clusters.

Selecting number of clusters

Here we'll compare the outputs of hierarchical clustering with the actual diagnoses. Because we have it in the dataset we can check the performance of our clustering model.

Use `cutree()` to cut the tree into 4 clusters, assign to `wisc.hclust.clusters`:

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=19)
```

Use the `table` function to compare to the actual diagnoses:

```
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
# 2 Clusters:  
wisc.hclust_2 <- cutree(wisc.hclust, h=25)  
table(wisc.hclust_2, diagnosis)
```

	diagnosis	
wisc.hclust_2	B	M
1	357	210
2	0	2

```
# 3 Clusters:  
wisc.hclust_3 <- cutree(wisc.hclust, h=22)  
table(wisc.hclust_3, diagnosis)
```

	diagnosis	
wisc.hclust_3	B	M
1	355	205
2	2	5
3	0	2

```
# 5 Clusters:
wisc.hclust_5 <- cutree(wisc.hclust, h=18)
table(wisc.hclust_5, diagnosis)
```

	diagnosis		
wisc.hclust_5	B	M	
1	12	165	
2	0	5	
3	343	40	
4	2	0	
5	0	2	

```
# 10 Clusters:
wisc.hclust_10 <- cutree(wisc.hclust, h=13)
table(wisc.hclust_10, diagnosis)
```

	diagnosis		
wisc.hclust_10	B	M	
1	12	86	
2	0	59	
3	0	3	
4	331	39	
5	0	20	
6	2	0	
7	12	0	
8	0	2	
9	0	2	
10	0	1	


Cutting into two or three clusters gives one cluster containing most of the results, with the other clusters only containing a few values. Increasing above four clusters causes you to start to lose the defined clusters containing either the majority of the B or M values.

Using different methods

There are different “*methods*” we can use to combine points during hierarchical clustering procedures. These include “single”, “complete”, “average”, and “ward.D2”.

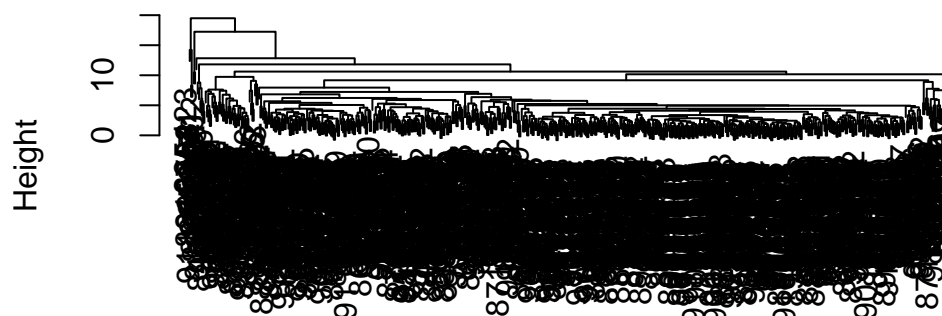
Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
wisc.hclust_single <- hclust(data_dist, method="single")  
plot(wisc.hclust_single)
```



```
wisc.hclust_average <- hclust(data_dist, method="average")  
plot(wisc.hclust_average)
```


Cluster Dendrogram



```
data_dist  
hclust (*, "average")
```

```
wisc.hclust_w.D2 <- hclust(data_dist, method="ward.D2")  
plot(wisc.hclust_w.D2)
```

Cluster Dendrogram



```
data_dist  
hclust (*, "ward.D2")
```


I like the ward.D2 combination, it leaves all of the individual values on the same level of clusters which I think makes the data look more clean.

K-means clustering

```
wisc.km <- kmeans(wisc.data, centers = 2)
table(wisc.km$cluster, diagnosis)
```

```
diagnosis
  B    M
1 356  82
2   1 130
```

Q14. How well does k-means separate the two diagnoses? How does it compare to hclust()?

It separates the benign diagnoses incredibly well, with only 1 ending up in a separate cluster. You lose a lot of data for the malignant diagnoses though. It is both better and worse compared to the hclust() model depending on how you look at it.

Use table() to compare k-means clusters with hierarchical clusters.

```
table(wisc.hclust.clusters, wisc.km$cluster)
```

```
wisc.hclust.clusters  1  2
                     1 68 109
                     2   5   2
                     3 365  18
                     4   0   2
```

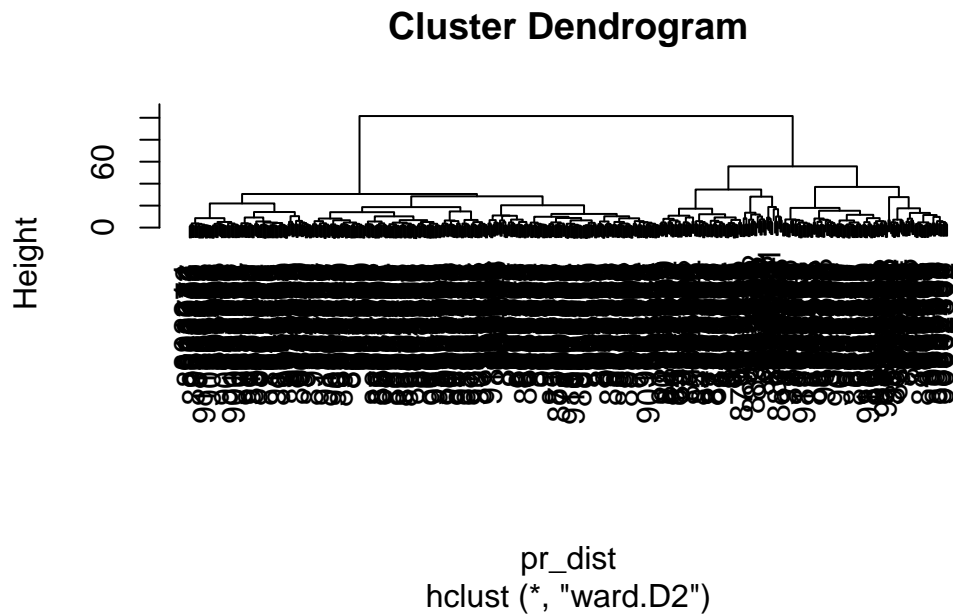
Combining Methods

Clustering on PCA results

Recall that the PCA required significantly fewer features to describe 70%, 80%, and 95% of the variability. PCA other has other benefits like normalizing data, avoiding over-fitting, and uncorrelating variables. Let's see if PCA improves or degrades the performance of hierarchical clustering.

Create a hierarchical clustering model using ward.D2 linkage to describe at least 90% of variability.

```
pr_dist <- dist(wisc.pr$x[,1:7])
pr.hclust <- hclust(pr_dist, method="ward.D2")
plot(pr.hclust)
```



This doesn't look anymore promising than our previous clustering models. Are the two main groups here malignant and benign?

```
grps <- cutree(pr.hclust, k=2)
table(grps)
```


```
grps
 1  2
216 353
```

```
table(grps, diagnosis)
```


```
diagnosis
```

```
grps  B  M
1  28 188
2 329  24
```

```
plot(wisc.pr$x, col=grps)
```



```
# Or:
plot(wisc.pr$x[,1:2], col=diagnosis)
```



Use the distance along the first 7 PCs for clustering:

```
wisc.pr.hclust <- hclust(pr_dist, method = "ward.D2")

# Cut this model into 2 clusters:
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

Using `table()` compare the results of the new hierarchical model with the actual diagnoses.

```
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

Q15. How well does the newly created model with four clusters separate the two diagnoses?

It works the same as the previous “grps” method. Still better than the original models done with the “complete”.

Q16. How well do the k-means and hierarchical clustering models do in terms of separating diagnoses? Use `table()` to compare the output of each model.

```
table(wisc.km$cluster, diagnosis)
```

```
diagnosis
  B    M
1 356  82
2   1 130
```

```
table(wisc.hclust.clusters, diagnosis)
```

```
           diagnosis
wisc.hclust.clusters  B    M
1      12 165
2       2   5
3     343  40
4       0   2
```

Both seem to struggle to really separate the benign diagnoses, though the `hclust()` model does a slightly better job at this. Even then the `hclust()` model does not have as good of a separation for benign diagnoses compared to the kmeans. Both have benefits and consequences.

Sensitivity and specificity

Sensitivity: a test's ability to correctly detect ill patients with the condition. In this case, the number of samples in the cluster identified as predominantly malignant divided by the total number of known malignant samples ($TP/(TP+FN)$)

Specificity: a test's ability to correctly reject healthy patients. In this case, the proportion of benign samples in the cluster identified as predominantly benign that are known to be benign ($TN/(TN+FP)$)

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? Sensitivity?

```
# Kmean sensitivity and specificity
(130/212)
```

```
[1] 0.6132075
```

```
(356/357)
```

```
[1] 0.9971989
```

```
# hclust sensitivity and specificity  
(165/212)
```

```
[1] 0.7783019
```

```
(343/357)
```

```
[1] 0.9607843
```

K-means produced the more specific but less sensitive model compared to hclust.

Prediction


We'll use `predict()` function that will take our PCA model from before and new cancer cell data and project that data onto our PCA space.

```
url <- "https://tinyurl.com/new-samples-CSV"  
new <- read.csv(url)  
npc <- predict(wisc.pr, newdata=new)  
head(npc)
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	

	PC27	PC28	PC29	PC30
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col = "white")
```



Q18. Which of these new patients should we prioritize for follow up based on the results?

Patients in group 2 should be prioritized for follow up appointments since they're more likely to have a malignant diagnosis.